

```
In [1]: import csv
import pandas as pd
import numpy as np
import scipy as sp
import collections
import math
import operator
import sklearn.metrics
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
```

```
In [2]: !pwd
!ls
```

```
/Users/xiasong/Documents/Class_2016/DSE/DSE220/homework/homework_2
Homework_2.pdf          mnist_test_data.csv    mnist_train_labels.csv
hw_2.ipynb              mnist_test_labels.csv
hw_2.pdf                mnist_train_data.csv
```

load data into python and check data structure

```
In [3]: train = pd.read_csv('mnist_train_data.csv')
train_lab = pd.read_csv('mnist_train_labels.csv')
test = pd.read_csv('mnist_test_data.csv')
test_lab = pd.read_csv('mnist_test_labels.csv')
```

```
In [4]: print (train.shape, train_lab.shape, test.shape, test_lab.shape)

(5999, 784) (5999, 1) (999, 784) (999, 1)
```

Question 1: Compute and report the prior probabilities π_j for all labels. (10 marks)

```
In [5]: train_lab['5'].head(2)
```

```
Out[5]: 0      0
        1      4
        Name: 5, dtype: int64
```

```
In [6]: #count each class in train lab
lab_count1=dict(collections.Counter(train_lab['5']))
lab_count = collections.OrderedDict(sorted(lab_count1.items()))
print (lab_count, lab_count1)
print ("-----")
#calculate the fraction of each digit
priors1=[]
for k, v in lab_count.items():
    a = (v/len(train_lab))
    a = round(a,4)
    priors1.append((k,a))
priors2 = dict(priors1)
priors = collections.OrderedDict(sorted(priors2.items()))
print ('The prior probabilities of  $\pi_j$  are:')
priors

OrderedDict([(0, 592), (1, 671), (2, 581), (3, 608), (4, 623), (5, 513),
              (6, 608), (7, 651), (8, 551), (9, 601)]) {0: 592, 4: 623, 1: 671, 9: 60
1, 2: 581, 3: 608, 5: 513, 6: 608, 7: 651, 8: 551}
-----
The prior probabilities of  $\pi_j$  are:
```

```
Out[6]: OrderedDict([(0, 0.0987),
                    (1, 0.1119),
                    (2, 0.0968),
                    (3, 0.1014),
                    (4, 0.1039),
                    (5, 0.0855),
                    (6, 0.1014),
                    (7, 0.1085),
                    (8, 0.0918),
                    (9, 0.1002)])
```

```
In [7]: lab_count1[1]
```

```
Out[7]: 671
```

Question 2: For each pixel X_i and label j , compute $P_{ji} = P(X_i = 1|y = j)$ (Use the maximum likelihood estimate shown in class). Use Laplacian Smoothing for computing P_{ji} . Report the highest P_{ji} for each label j . (15 marks)

```
In [8]: priors[1]
```

```
Out[8]: 0.1119
```

```
In [9]: #calcualte the pji
df1 = pd.concat([train_lab, train], axis=1)
Pji = []
for j in range(10):
    for i in range(1,785):
        df2 = df1.ix[:,(0,i)] #from df1 extract the 0 and ith column
        df3 = df2[(df2.iloc[:,0] == j) & (df2.iloc[:,1] > 0.9)]
        pji = (len(df3)+1) / (lab_count1[j]+2)
        Pji.append((j,pji))
```

```
In [10]: #select the highest Pji for each label
k = 784
Pj = []
for j in range(10):
    n = []
    for i in range (784):
        a = j * k + i
        b = Pji[a][1]
        b = round(b,4)
        n.append(b)
    c = max(n)
    Pj.append((Pji[a][0],c))
#Pj = dict(Pj)
print ("The highest Pji for each label j are")
Pj
```

The highest Pji for each label j are

```
Out[10]: [(0, 0.8519),
(1, 0.9851),
(2, 0.729),
(3, 0.8082),
(4, 0.8496),
(5, 0.7126),
(6, 0.8492),
(7, 0.7948),
(8, 0.8752),
(9, 0.8673)]
```

Question 3: Use naive bayes (as shown in lecture slides) to classify the test data. Report the accuracy. (5 marks)

```
In [11]: # form the classifier
k = 784
argmax = []
for i in range(len(test)):
    a = test.iloc[i,:]
    am = []
    for j in range(10):
        am1 = []
        for m in range(784):
            d = k * j + m
            if a[m] > 0.9:
                am2 = math.log(Pji[d][1])
                am1.append(am2)
        am3 = math.log(priors[j]) + sum(am1)
        am.append(am3)
    idx, val = max(enumerate(am), key=operator.itemgetter(1))
    argmax.append((idx, val))
```

```
In [12]: # evaluate the model results with accuracy
labels_pre = []
for i in range(len(argmax)):
    a = argmax[i][0]
    labels_pre.append(a)
count = 0
for i in range(len(labels_pre)):
    if labels_pre[i] == test_lab.iloc[i,0]:
        count = count + 1
accuracy = count / len(test)

print ('The accuracy of model is %s' % accuracy)
```

The accuracy of model is 0.6706706706706707

Question 4: Compute the confusion matrix (as shown in the lectures) and report the top 3 pairs with most (absolute number) incorrect classifications. (10 marks)

```
In [13]: labels_tru = []
for i in range(len(test_lab)):
    a = test_lab.iloc[i,0]
    labels_tru.append(a)
```

```
In [14]: confusion_matrix(labels_tru, labels_pre)
```

```
Out[14]: array([[81,  0,  0,  0,  0,  0,  0,  0,  4,  0],
 [ 0, 60,  1,  2,  0,  0,  1,  0, 62,  0],
 [ 5,  0, 92,  0,  0,  0,  1,  0, 18,  0],
 [ 0,  0,  2, 91,  0,  1,  3,  0,  9,  1],
 [ 2,  0,  2,  0, 59,  0,  3,  0, 28, 16],
 [12,  0,  2, 18,  1, 10,  2,  1, 40,  1],
 [ 8,  0,  4,  0,  1,  0, 68,  0,  6,  0],
 [ 3,  0,  2,  3,  1,  0,  0, 58, 21, 10],
 [ 3,  0,  2,  8,  0,  0,  0,  0, 76,  0],
 [ 1,  0,  1,  2,  2,  0,  0,  0, 13, 75]])
```


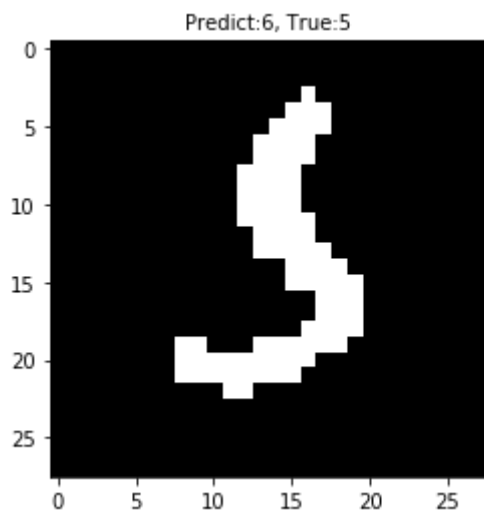
From the confusion_matrix we can tell that the top 3 pairs of incorrect classification are (8,1), (8,5) and (8,4).

Question 5: Visualizing mistakes: Print two MNIST images from the test data that your classifier misclassified. Write both the true and predicted labels for both of these misclassified digits. (10 marks)

```
In [15]: #find the index where the predicted labels do not match with true labels
indicies = []
for i in range(len(labels_tru)):
    if labels_tru[i] != labels_pre[i]:
        indicies.append(i)
```

```
In [16]: pixels1 = np.array(test.iloc[339,:], dtype='uint8')
pixels1 = np.array(pixels1, dtype='uint8')
pixels2 = np.array(test.iloc[267,:], dtype='uint8')
pixels2 = np.array(pixels2, dtype='uint8')
# Reshape the array into 28 x 28 array (2-dimensional array)
pixels1 = pixels1.reshape((28, 28))
pixels2 = pixels2.reshape((28, 28))
plt.title('Predict:%i, True:%i' % ( labels_pre[339],labels_tru[339]), fontsize=12)
plt.imshow(pixels1, cmap='gray')
plt.show()

plt.title('Predict:%i, True:%i' % ( labels_pre[267],labels_tru[267]), fontsize=12)
plt.imshow(pixels2, cmap='gray')
plt.show()
```



Now, we will implement Gaussian Mixture Model and Linear Discriminant Analysis on the breast cancer data (sklearn.datasets.load_breast_cancer) available in sklearn.datasets. Load the data and split it into train-validation-test (40-20-40 split). Don't shuffle the data, otherwise your results will be different.

```
In [39]: #download data
from sklearn import datasets
cancer = datasets.load_breast_cancer()
X = cancer.data
y = cancer.target
type(y)
```

```
Out[39]: numpy.ndarray
```

```
In [40]: X_tra_val, X_tes, y_tra_val, y_tes = train_test_split(X, y, test_size=0.4, r
X_tra, X_val, y_tra, y_val = train_test_split(X_tra_val, y_tra_val, test_size
```

Question 6: Implement Gaussian Mixture model on the data as shown in class. Tune the covariance type parameter on the validation data. Use the selected value to compute the test accuracy. As always, train the model on train+validation data to compute the test accuracy. (10 mark)

```
In [41]: from sklearn.mixture import GaussianMixture
from sklearn.metrics import accuracy_score

# Initialize Gaussian Naive Bayes
for cov_type in ('full', 'tied', 'diag', 'spherical'):
    gm = GaussianMixture(n_components=3, covariance_type = cov_type)
    # Train the classifier
    gm.fit(X_val, y_val)
    # Make predictions on test data
    y_pre = gm.predict(X_tes)
    accuracy = accuracy_score(y_pre, y_tes)
    print ('Test accuracy = ' + str(accuracy) + ' at covariance_tpye = ' + c
```

```
Test accuracy = 0.175438596491 at covariance_tpye = full
Test accuracy = 0.0614035087719 at covariance_tpye = tied
Test accuracy = 0.859649122807 at covariance_tpye = diag
Test accuracy = 0.179824561404 at covariance_tpye = spherical
```

```
In [43]: #train the model on train+validation compute the test accuracy
gm = GaussianMixture(n_components=3, covariance_type = 'diag')
# Train the classifier
gm.fit(X_tra_val, y_tra_val)
# Make predictions on test data
y_pre = gm.predict(X_tes)
accuracy = accuracy_score(y_pre, y_tes)
print ('Test accuracy = ' + str(accuracy))
```

```
Test accuracy = 0.736842105263
```

Question 7: Apply Linear Discriminant Analysis model on the train+validation data and report the accuracy obtained on test data. Report the transformation matrix (w) along with the intercept. (5 mark)

```
In [45]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
         clf = LinearDiscriminantAnalysis()
         # Train
         clf.fit(X_tra_val, y_tra_val)
         # Test
         y_pre = clf.predict(X_tes)

         # print the accuracy
         print ('Test accuracy = ' + str(np.sum(y_pre == y_tes)/len(y_tes)))

Test accuracy = 0.964912280702
```

```
In [46]: clf.transform(X_tra_val)
```

```
Out[46]: array([[ 2.26351958],  
                [ 1.50315616],  
                [-2.32016749],  
                [ 1.16950535],  
                [ 0.2848613 ],  
                [-1.64129414],  
                [-4.67284807],  
                [ 1.07672183],  
                [ 1.54651486],  
                [ 0.64701382],  
                [ 2.6571233 ],  
                [ 2.6550835 ],  
                [ 2.33859685],  
                [-1.94301274],  
                [-0.62044859],  
                [ 1.56540432],  
                [-2.78728741],  
                [-1.88634026],  
                [-3.01128946],  
                [ 2.01215895],
```

Question 8: Load the digits dataset (scikit-learn's toy dataset) and take the last 1300 samples as your test set. Train a K-Nearest Neighbor ($k=5$, linf distance) model and then without using any scikit-learn method, report the nal values for Specicity, Sensitivity, TPR, TNR, FNR, FPR, Precision and Recall for Digit 3 (this digit is a positive, everything else is a negative). (15 marks)

```
In [2]: from sklearn import datasets
digits= datasets.load_digits()
X = digits.data
y = digits.target
```

```
In [3]: len(X)
```

Out[3]: 1797

```
In [4]: #split dataset into train and test
X_tra = X[0:497]
X_tes = X[498:]
y_tra = y[0:497]
y_tes = y[498:]
```

```
In [5]: #train the model in f classes
from sklearn.neighbors import KNeighborsClassifier
clf = KNeighborsClassifier(5)
clf.fit(X_tra, y_tra)
#predict test data
pred = clf.predict(X_tes)
```

```
In [6]: #Calculate TPR, Sensitivity and Recall
TP = 0
FN = 0
for i in range(len(y_tes)):
    if ((y_tes[i] == 3) & (pred[i] == 3)):
        TP = TP + 1
    elif ((y_tes[i] == 3) & (pred[i] != 3)):
        FN = FN + 1
```

```
In [7]: #Calculate TPR, Sensitivity and Recall
TPR = TP / (TP + FN)
TPR
```

Out[7]: 0.8769230769230769

```
In [9]: #Calculate TNR, Specificity
TN = 0
FP = 0
for i in range(len(y_tes)):
    if ((y_tes[i] != 3) & (pred[i] != 3) & (y_tes[i] == pred[i])):
        TN = TN + 1
    elif ((y_tes[i] != 3) & (pred[i] == 3)):
        FP = FP + 1
```

```
In [10]: #Calculate TNR, Specificity
TNR = TN / (TN + FP)
TNR
```

Out[10]: 0.9880294659300184

```
In [11]: #Calculate FNR
FNR = FN / (TP + FN)
FNR
```

Out[11]: 0.12307692307692308

```
In [12]: #Calculate FPR
FPR = FP / (FP + TN)
FPR
```

Out[12]: 0.011970534069981584

```
In [13]: #Calculate precisions
Precisions = TP / (TP + FP)
Precisions
```

Out[13]: 0.8976377952755905

An ablation experiment consists of removing one feature from an experiment, in order to assess the amount of additional information that feature provides above and beyond the others. For this section, we will use the diabetes dataset from scikit-learn's toy datasets. Split the data into training and testing data as a 90-10 split with random state of 10.

```
In [58]: from sklearn import datasets
diabetes= datasets.load_diabetes()
X = diabetes.data
y = diabetes.target
```

```
In [59]: X_tra, X_tes, y_tra, y_tes = train_test_split(X, y, test_size=0.1, random_st
```

Question 9: Perform least squares regression on this dataset. Report the mean squared error and the mean absolute error on the test data. (5 marks)

```
In [60]: from sklearn.linear_model import LinearRegression
```

```
In [61]: # Least squares regression
theta,residuals,rank,s = np.linalg.lstsq(X_tra, y_tra)
# Make predictions on the test data
predictions = np.dot(X_tes, theta)
```

```
In [14]: from sklearn.linear_model import LinearRegression
lin = LinearRegression()
lin.fit(X_tra, y_tra)
predictions = lin.predict(X_tes)
```

```
/Users/xiasong/anaconda2/envs/py36/lib/python3.6/site-packages/scipy/linalg/basic.py:1018: RuntimeWarning: internal gelsd driver lwork query error, required iwork dimension not returned. This is likely the result of LAPACK bug 0038, fixed in LAPACK 3.2.2 (released July 21, 2010). Falling back to 'gelss' driver.
warnings.warn(msg, RuntimeWarning)
```

```
In [63]: # Mean squared error calculation
from sklearn.metrics import mean_squared_error
print (mean_squared_error(y_tes, predictions))

2155.96465103
```

```
In [64]: # Mean absolute error calculation
from sklearn.metrics import mean_absolute_error
print (mean_absolute_error(y_tes, predictions))

36.3181336987
```

Question 10: Repeat the experiment from Question 10 for all possible values of ablation (i.e., removing the feature 1 only, then removing the feature 2 only, and so on). Report all MSEs. (10 marks)

```
In [65]: len(X_tra)
```

```
Out[65]: 397
```

```
In [66]: a = np.delete(X_tra, np.s_[1],1)
len(a)
```

```
Out[66]: 397
```

```
In [67]: for i in range(10):
X_tra_rem = np.delete(X_tra, np.s_[i],1)
X_tes_rem = np.delete(X_tes, np.s_[i],1)
lin = LinearRegression()
lin.fit(X_tra_rem, y_tra)
predictions = lin.predict(X_tes_rem)
a = mean_squared_error(y_tes, predictions)
print ('MSE = ' + str(a) + ' when remove the %i feature' %i)
```

```
MSE = 2152.80664218 when remove the 0 feature
MSE = 2259.13307937 when remove the 1 feature
MSE = 2783.51448185 when remove the 2 feature
MSE = 2424.772348 when remove the 3 feature
MSE = 2187.59951938 when remove the 4 feature
MSE = 2167.51760615 when remove the 5 feature
MSE = 2159.15148251 when remove the 6 feature
MSE = 2153.06317113 when remove the 7 feature
MSE = 2335.17338461 when remove the 8 feature
MSE = 2165.86619219 when remove the 9 feature
```

Question 11: Based on the MSE values obtained from Question 11, which features do you deem the most/least significant and why? (5 marks)

According to the results of questions 10, we can draw the conclusion that the second is the most significant and the 7th feature is the least significant feature. Actually, we can tell the importance of each feature from the ESMs change before and after removing these features. When we remove the 2nd feature, the magnitude change of MSE is the largest, and when we remove the 7th feature, the magnitude change of MSE is least. Therefore, the second feature is most significant feature as to this linear regression and the 7th feature is the least significant feature to this linear regression.

```
In [ ]:
```